

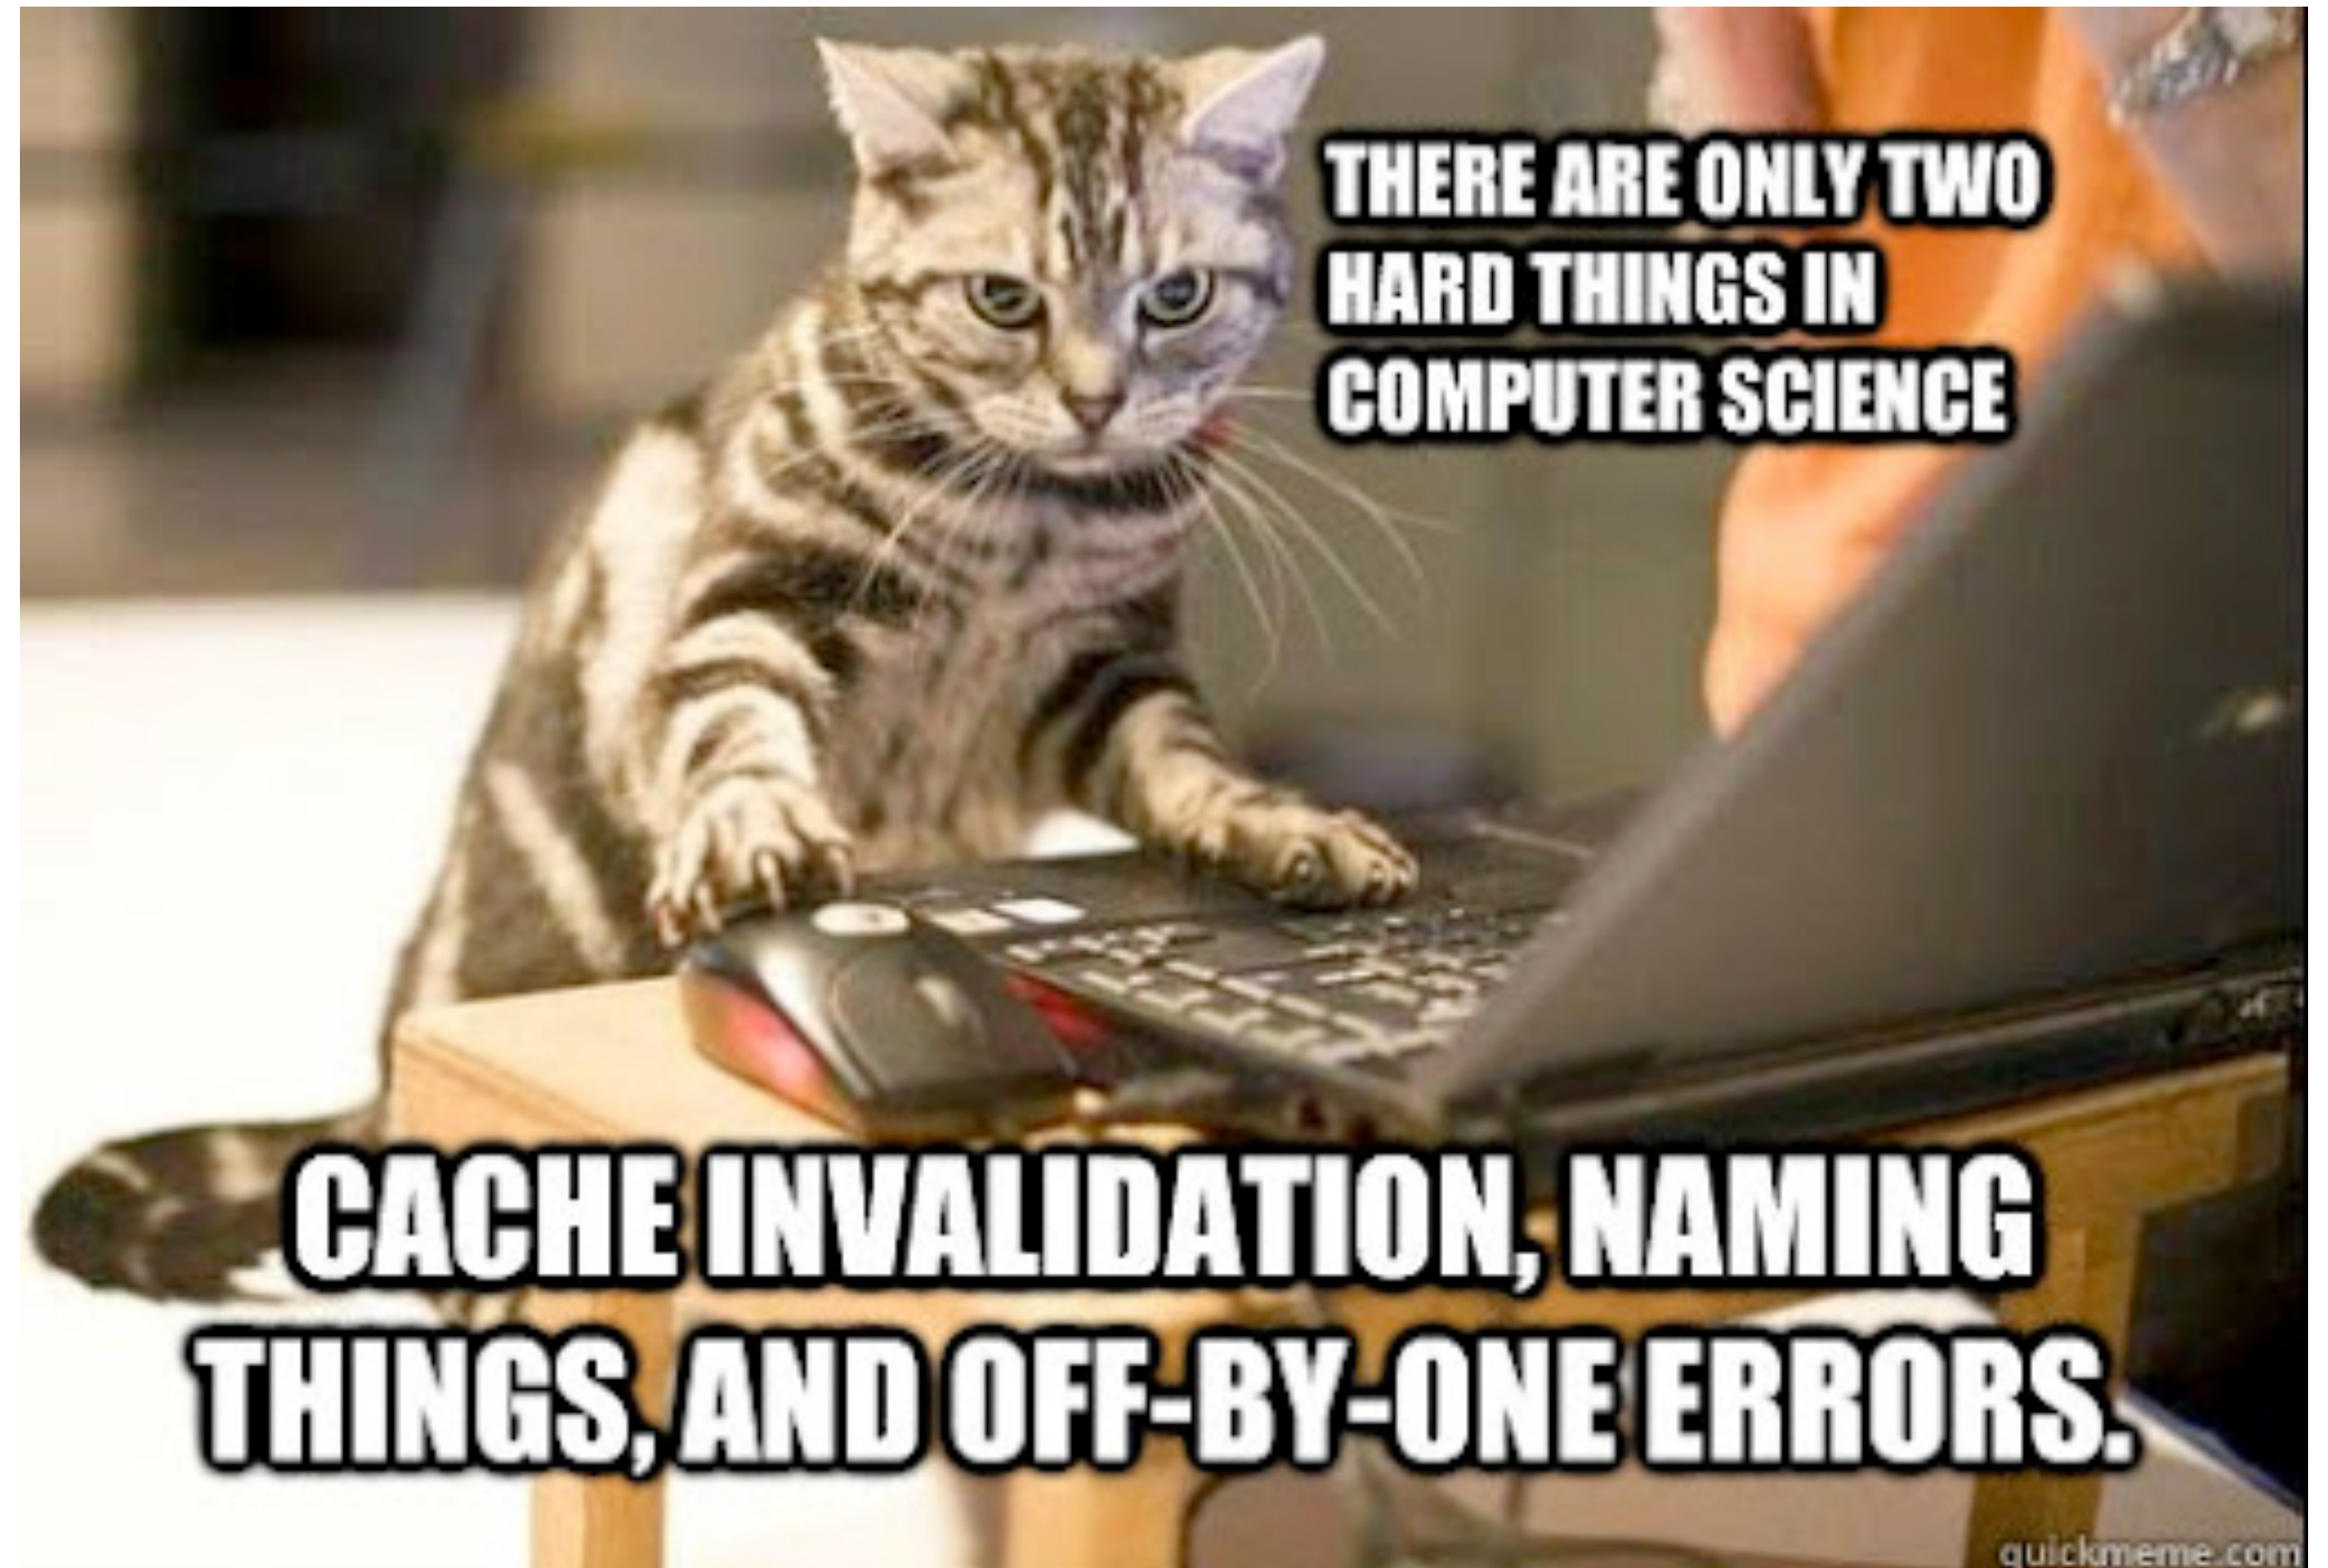
Data Cache & Locality

Based on slides by David Merchant and Troels Henriksen

Axel Kanne

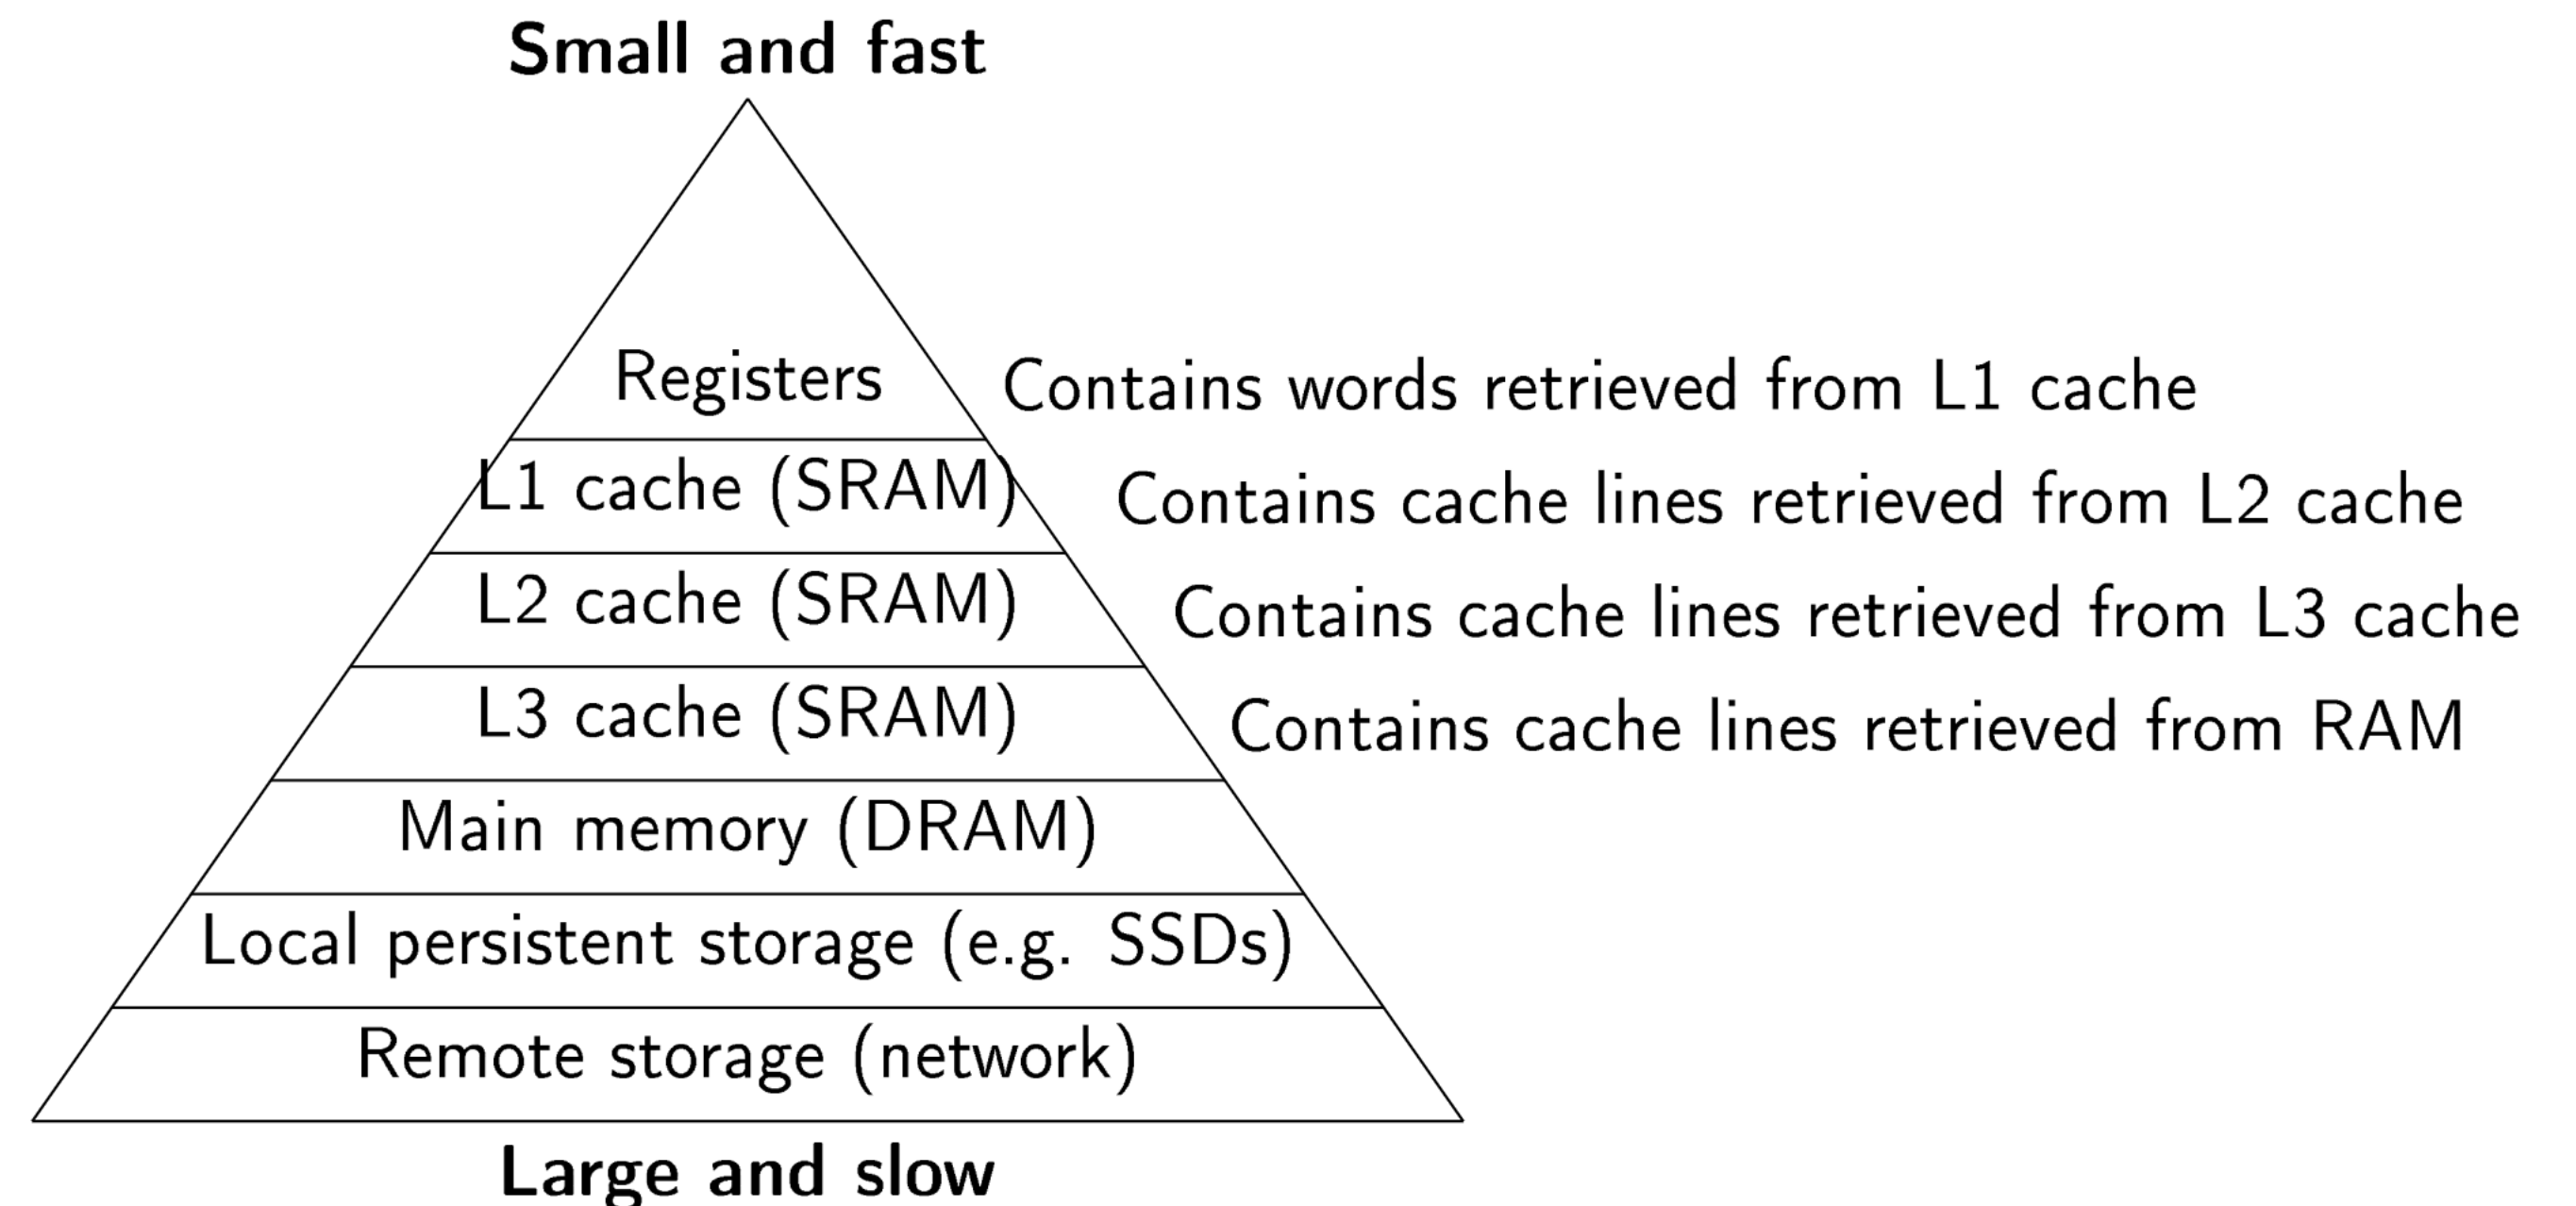
Agenda

- Theory
- Locality of code
- Cache exercises



The memory hierarchy

- Small memory: fast but small (it requires more power, and each storage unit is expensive)
- Large memory: slow but large (requires less power, and each storage unit is cheap)
- Solution: The smaller and faster device at level k acts as a cache for the larger slower device at level $k + 1$.
- **Cache:** A smaller, faster storage device that acts as a staging area for a subset of the data in a larger, slower device.



Borrowed from slides: Memory Hierarchy and Caching

Array layouts: row-major and column-major

- **Row-major order:** rows are contiguous in memory
 - Used by C
- **Column-major order:** columns are contiguous:
 - Used by MATLAB

Example: $A[3][3] = \{\{1, 2, 3\}, \{4, 5, 6\}, \{7, 8, 9\}\}$

Row-major:

Row 1			Row 2			Row 3		
1	2	3	4	5	6	7	8	9

Column-major:


Column 1			Column 2			Column 3		
1	4	7	2	5	8	3	6	9

Example continued: row major

Example:

```
int A[3][3] =  
    {{1, 2, 3},  
     {4, 5, 6},  
     {7, 8, 9}}
```

- Assume integers in the array are 4 bytes
- Assume we have a cache with blocks of size 12 ($\approx 3 * \text{sizeof}(\text{int})$), that can only hold 1 block at a time.
- What happens to a cold cache (a cache with nothing in it) where we access the row and second column of our 2D array, assuming row-major order is used? ($A[0][1]$)




Example continued: row major

Example:

```
int A[3][3] =  
    {{1,2,3},  
     {4,5,6},  
     {7,8,9}}
```

- Assume integers in the array are 4 bytes
- Assume we have a cache with blocks of size 12 ($\approx 3 \times \text{sizeof}(\text{int})$), that can only hold 1 block at a time.
- What happens to a cold cache (a cache with nothing in it) where we access the row and second column of our 2D array, assuming row-major order is used? ($A[0][1]$)




Example continued: row major

Example:

```
int A[3][3] =  
    {{1, 2, 3},  
     {4, 5, 6},  
     {7, 8, 9}}
```

- Assume integers in the array are 4 bytes
- Assume we have a cache with blocks of size 12 ($\approx 3 * \text{sizeof}(\text{int})$), that can only hold 1 block at a time.
- What happens to a cold cache (a cache with nothing in it) where we access the row and second column of our 2D array, assuming row-major order is used? ($A[0][1]$)

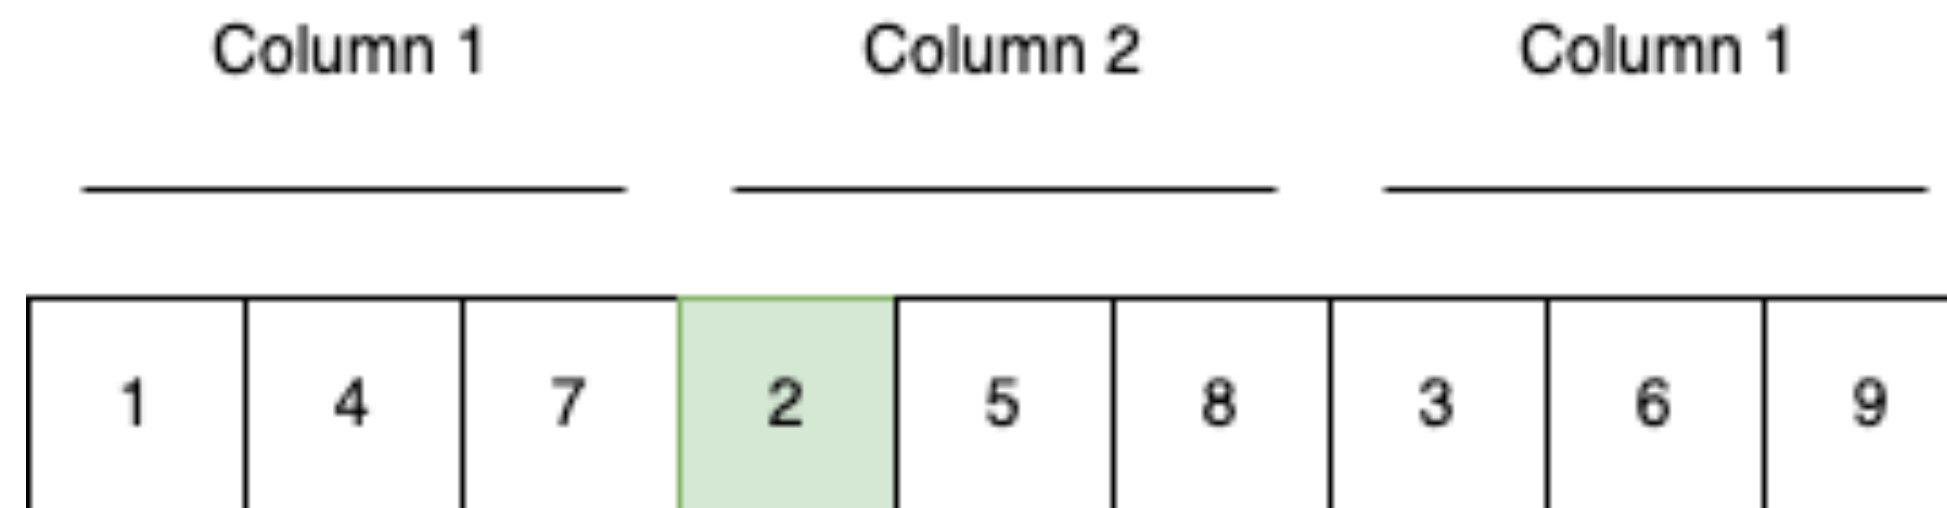


Example continued: column-major

Example:

```
int A[3][3] =  
    {{1,2,3},  
     {4,5,6},  
     {7,8,9}}
```

- Assumptions are the same as the last slide
- What happens to a cold cache where we access the first row and second column of our 2D array, assuming column-major order is used? ($A[0][1]$)

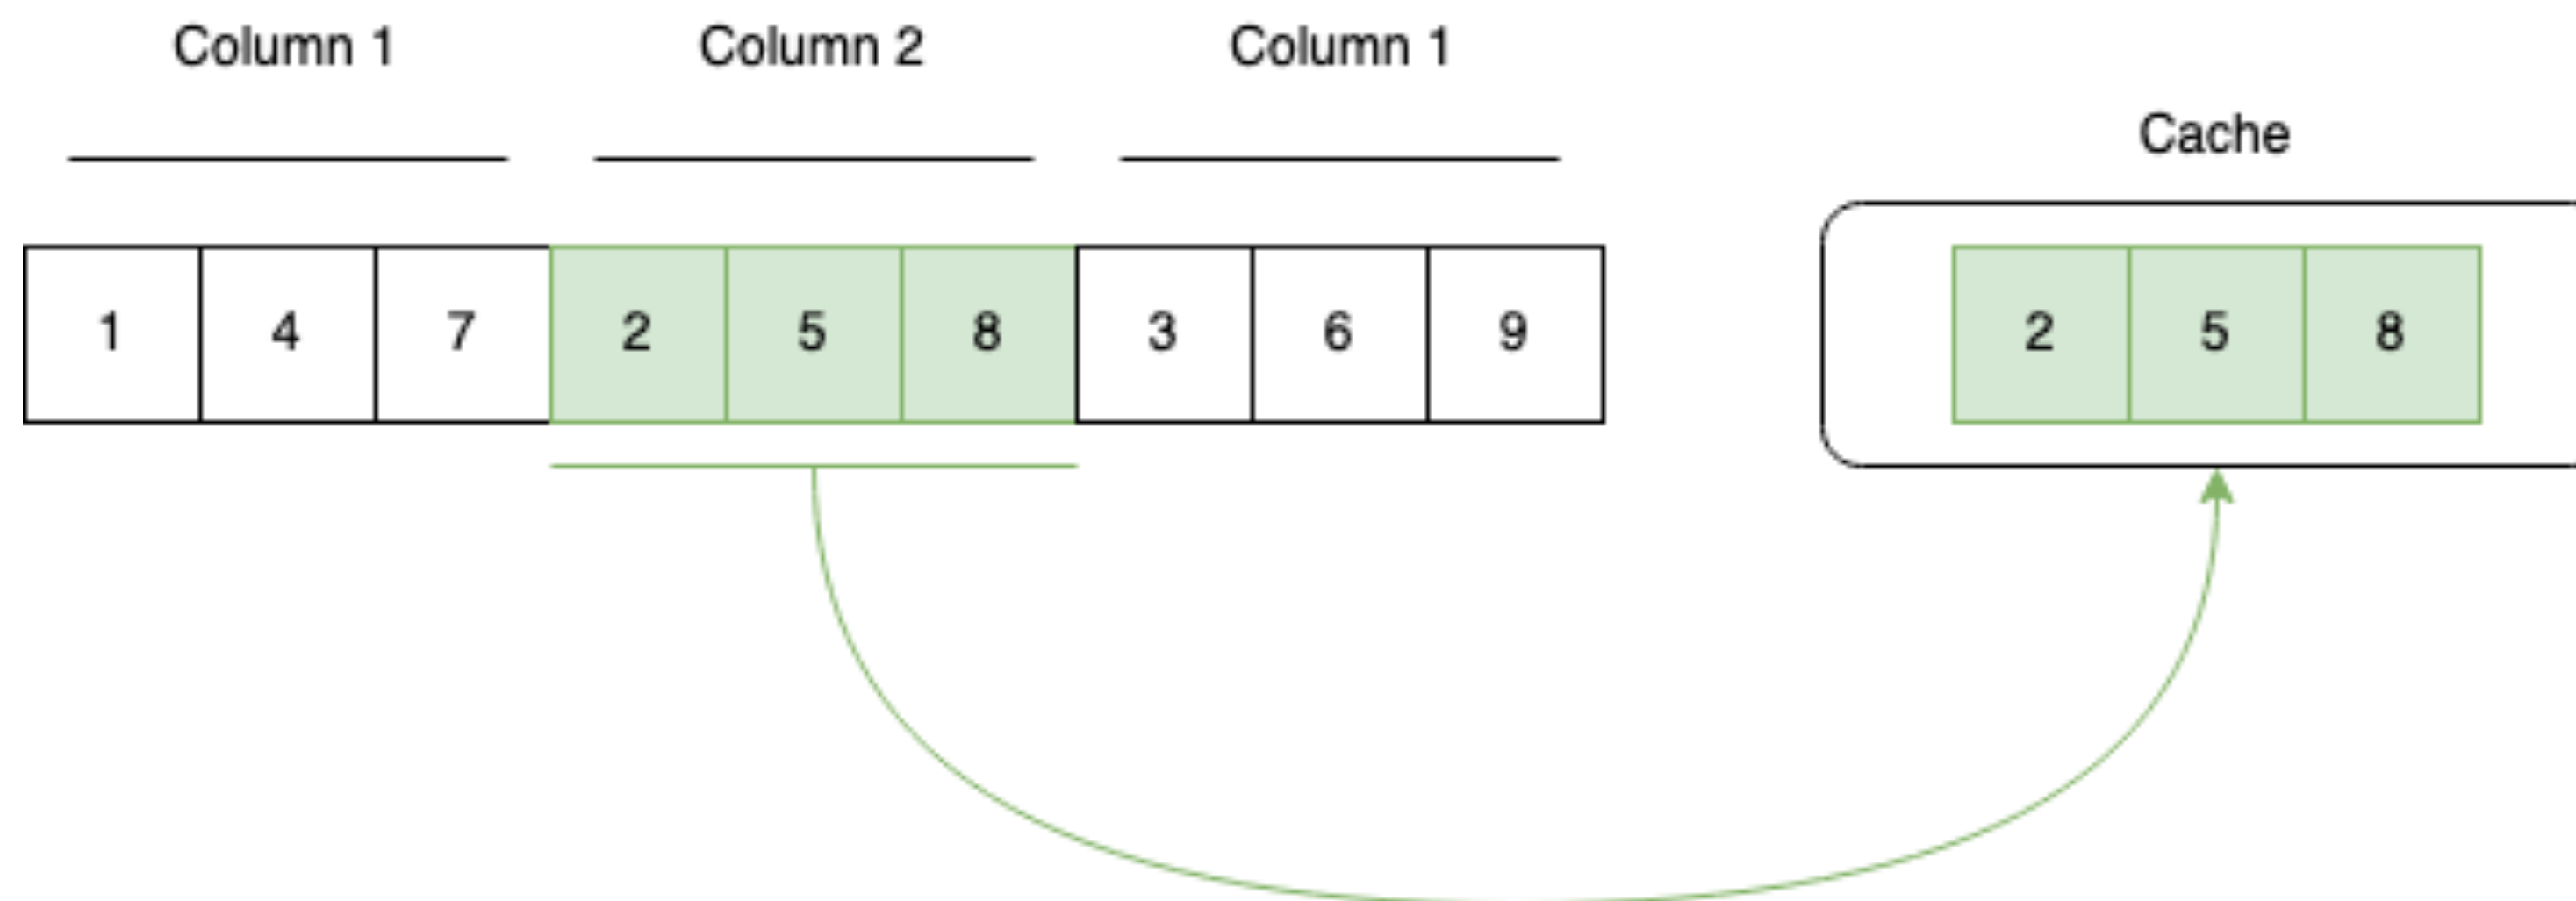


Example continued: column-major

Example:

```
int A[3][3] =  
    {{1,2,3},  
     {4,5,6},  
     {7,8,9}}
```

- Assumptions are the same as the last slide
- What happens to a cold cache where we access the first row and second column of our 2D array, assuming column-major order is used? ($A[0][1]$)



Locality

- **Spatial locality:** Data tends to be accessed at the same place as recently accessed memory.
- **Temporal locality:** Data that has recently been accessed, tends to be accessed again soon.
- **Stride:** How large are the jumps between memory accesses?

Locality: Example 1

```
1  int A[3][3] = {
2      {1, 2, 3},
3      {4, 5, 6},
4      {6, 7, 8}
5  };
6  int sum = 0;
7  for (int i = 0; i < 3; i++) {
8      for (int j = 0; j < 3; j++) {
9          sum += A[i][j];
10     }
11 }
12 printf("%d", sum);
```

- Stride:
 - row-major = 1
 - column-major = 3
- Spatial locality:
 - row-major = Better (we make jumps of 1 elements)
 - column-major = Worse (we make jumps of 3 elements)
- Temporal locality: `sum`, `i` and `j`


Locality: Example 2

```
1  ✓ int A[3][2][3] = {
2  ✓   {
3      {1, 2, 3}, {4, 5, 6}
4      },
5  ✓   {
6      {7, 8, 9}, {10, 11, 12}
7      },
8  ✓   {
9      {13, 14, 15}, {16, 17, 18}
10     }
11 };
12 int sum = 0;
13 ✓ for (int i = 0; i < 3; i++) {
14 ✓   for (int j = 0; j < 3; j++) {
15 ✓     for (int z = 0; z < 2; z++) {
16         sum += A[j][z][i];
17     }
18 }
19 }
20 printf("%d", sum);
```

- Is this the most optimal access order for either row-major or column-major?
 - Row-major: No, *i* that is used to access elements in the first dimension is incremented last
 - Column-major: No, *j*, that accesses 2-D arrays in the third dimension is not incremented first.
- What is the most optimal arrangement of the for-loops (from outer to inner)?
 - Row-major: *j*, *z*, *i*
 - Column-major: *i*, *z*, *j*

Cache organisation

- **N-way set associative:**
 - N: how many blocks there are in each set.
- **Direct mapped/One-way set associative:** one block for each set, so each memory location is mapped to one location in the cache.
- **Fully associative:** cache is one set, a memory location can be mapped to anywhere in the cache.
 - Example: If the cache has space for 8 blocks, and it is fully associative, it would be an 8-way set associative.



Borrowed from slides: Memory Hierarchy and Caching

Address translation

- Order in bits left to right: tag, set index, block offset.
- Offset bits (b): how many bits do you need to represent a block?

$$b = \log_2(\text{blocksize_in_bytes})$$

- How many sets are in an n-way associative cache?


$$\text{set_count} = \frac{\text{cachesize_in_bytes}}{n \cdot \text{blocksize_in_bytes}}$$

- Set index bits (s): how many bits do you need to represent the sets?

$$s = \log_2(\text{set_count}) = \log_2\left(\frac{\text{cachesize_in_bytes}}{n \cdot \text{blocksize_in_bytes}}\right)$$

- Tag: the rest of the address: $t = \text{bits_in_address} - (s + b)$

Address components



Borrowed from slides: Memory Hierarchy and Caching

Unit conversion

- Remember to convert between the correct units:
 - Kibibyte (KiB) = 1024 bytes
 - Mebibyte (MiB) = 1024 KiB
 - kb = kbit = kilobit = 125 bytes
- Memory is byte-addressed, so it is a good idea to convert everything to bytes.

Exam 2022-2023 question 1.1.1

1.1 Data Cache and Locality (about 8 %)

A byte-addressed machine with $\text{bits_in_address} = 32$ is equipped with a $\text{cachesize} = 16\text{KiB}$ $n = 4$ way set-associative cache with a $\text{blocksize} = 32$ bytes. The cache uses “least-recently-used replacement policy”.

Question 1.1.1: The address is separated into the following fragments when the cache is accessed

- block offset,
- cache tag, and
- set index.

$$\begin{aligned} b &= \log_2(\text{blocksize_in_bytes}) = \log_2(32) = 5 \\ \text{cachesize_in_bytes} &= 16 \text{ KiB} * 1024 = 16384 \text{ bytes} \\ \text{set_count} &= \frac{\text{cachesize_in_bytes}}{n \cdot \text{blocksize_in_bytes}} = \frac{16384}{4 \cdot 32} = 128 \\ s &= \log_2(\text{set_count}) = \log_2(128) = 7 \\ t &= \text{bits_in_address} - (s + b) = 32 - (7 + 5) = 20 \end{aligned}$$

What is bit-size of each and how are they ordered?

	31		0
Address fragment			
	Cache tag	Set index	Block offset
Bit-size	20	7	5



Exam 2022-2023 question 1.1.2

	31		0
Address fragment	Cache tag	Set index	Block offset
Bit-size	20	7	5

Question 1.1.2: Calculate for the following addresses the value of block offset, cache tag, and set index.
Address: 0x76543210

Cache tag = 0x76543

Block offset = 0x10

0x76543210 = 0111 0110 0101 0100 0011 0010 0001 0000

Set index = 0x10

Replacement schemes

- **LRU (least recently used)**: the block that is replaced in the set is the one that has been unused for the longest time.
- **Random**: candidate blocks are randomly selected, possibly using hardware assistance.
- Assuming LRU:
 - Find the set index of the address.
 - Search for the block in the set by comparing tags.
 - If it is found = **cache hit**: update this block such that it is visible that it was just used.
 - On **cache miss**:
 - If there are one or more blocks that are not valid (check the valid bit), replace one of these blocks with the new block.
 - If all blocks in the set are valid, replace the least recently used block, and update the new block such that it is visible that it was just used.

Exam 2022-2023 question 1.1.3

1.1 Data Cache and Locality (about 8 %)

A byte-addressed machine with $\text{bits_in_address} = 32$ is equipped with a 16 kibibyte 4-way set-associative cache with a $\text{blocksize} = 32$ bytes. The cache uses "least-recently-used replacement policy".

Question 1.1.1: The address is separated into the following fragments when the cache is accessed

- block offset,
- cache tag, and
- set index.

What is bit-size of each and how are they ordered?

	31		0
Address fragment	Cache tag	Set index	Block offset
Bit-size	20	7	5

$\text{cachesize} = 16\text{KiB}$ $n = 4$

$b = \log_2(\text{blocksize_in_bytes}) = \log_2(32) = 5$

$\text{cachesize_in_bytes} = 16 \text{ KiB} * 1024 = 16384 \text{ bytes}$

$\text{set_count} = \frac{\text{cachesize_in_bytes}}{n \cdot \text{blocksize_in_bytes}} = \frac{16384}{4 \cdot 32} = 128$

$s = \log_2(\text{set_count}) = \log_2(128) = 7$

$t = \text{bits_in_address} - (s + b) = 32 - (7 + 5) = 20$



Exam 2022-2023 question 1.1.3

Question 1.1.3: On the machine, the following stream of cache accesses are performed (read from top to bottom). Indicate for each of the address references:

- the set index,
- if the cache access is a miss or a hit, and Only write tags of the affected set
- the cache tags in LRU-order that are in the affected set after the access.

Addresses are given in hexa-decimal notation. Assume the cache is cold on entry.
Cache is cold

Reference: 0x0100004 = 0000 0001 0000 0000 0000 0000 0100

Tag = 0x100

Set index = 0

Cache is cold (empty), so it is a miss.

Cache { 0: {0x100} }

Reference	Set index	Hit/Miss	State of Tags
0x0100004	0	miss	0x100

Exam 2022-2023 question 1.1.3

Cache before { 0: {0x100} }

Reference: 0x0100010 = 0000 0001 0000 0000 0000 0001 0000

Tag = 0x100 Set index = 0

Cache hit, state of cache does not change:

Cache after { 0: {0x100} }

0x0100010	0	hit	0x100
-----------	---	-----	-------

Exam 2022-2023 question 1.1.3

Cache before { 0: {0x100} }

Reference: 0x0010008 = 0000 0000 0001 0000 0000 0000 1000

Tag = 0x10

Set index = 0

Cache miss, set 0 is not full, so tag 0x10 is added to set 0, and 0x10 is first in LRU-order:

Cache after { 0: {0x10, 0x100} }

Important tip!

We write the tags in the set in LRU-order, such that it easy to keep track

Exam 2022-2023 question 1.1.3

Cache before { 0: {0x100} }

Reference: 0x0010008 = 0000 0000 0001 0000 0000 0000 1000

Tag = 0x10 Set index = 0

Cache miss, set 0 is not full, so tag 0x10 is added to set 0, and 0x10 is first in LRU-order:

Cache after { 0: {0x10, 0x100} }	0x0010008	0	miss	0x10,0x100
----------------------------------	-----------	---	------	------------

Exam 2022-2023 question 1.1.3

Cache before: { 0: {0x10, 0x100} }

Reference: 0x0110004 = 0000 0001 0001 0000 0000 0000 0100

Tag = 0x110 Set index =0

Cache miss, set 0 is not full, so tag 0x110 is added to set 0, and LRU-order is changed:

Cache after:

{ 0: {0x110, 0x10, 0x100} }

0x0110004	0	miss	0x110,0x10,0x100
-----------	---	------	------------------

Exam 2022-2023 question 1.1.3

Cache before: { 0: {0x110, 0x10, 0x100} }

Reference: 0x00000008 = 0000 0000 0000 0000 0000 0000 1000

Tag = 0x0

Set index = 0

Cache miss, set 0 is not full, so tag 0x0 is added to set 0, and LRU-order is changed:

Cache after:

{ 0: {0x0, 0x110, 0x10, 0x100} }

0x00000008	0	miss	0x0,0x110,0x10,0x100
------------	---	------	----------------------

Exam 2022-2023 question 1.1.3

Cache before: { 0: {0x0, 0x110, 0x10, 0x100} }

Reference: 0x0300004 = 0000 0011 0000 0000 0000 0000 0100

Tag = 0x300 Set index = 0

Cache miss, set 0 is full (it can only hold 4 blocks), so tag 0x100 is replaced with 0x300 since it is last in the LRU-order list.

Cache after:

{ 0: {0x300, 0x0, 0x110, 0x10} }

0x0300004	0	miss	
			0x300,0x0,0x110,0x10

Exam 2022-2023 question 1.1.3

Cache before: { 0: {0x300, 0x0, 0x110, 0x10} }

Reference: 0x0300010 = 0000 0011 0000 0000 | 0000 0001 0000

Tag = 0x300 Set index = 0

Cache hit, but tag 0x300 is already the first in the LRU-order list, so nothing is changed:

Cache after:

{ 0: {0x300, 0x0, 0x110, 0x10} }

0x0300010	0	hit	0x300,0x0,0x110,0x10
-----------	---	-----	----------------------

Exam 2022-2023 question 1.1.3

Cache before: { 0: {0x300, 0x0, 0x110, 0x10} }

Reference: 0x0110008 = 0000 0001 0001 0000 0000 0000 1000

Tag = 0x110 Set index = 0

Cache hit, LRU-order is updated in set 0 so tag 0x110 is first:

Cache after:

{ 0: {0x110, 0x300, 0x0, 0x10} }

0x0110008	0	hit	
			0x110,0x300,0x0,0x10

Exam 2022-2023 question 1.1.3

Cache before: { 0: {0x110, 0x300, 0x0, 0x10} }

Reference: 0x0410004 = 0000 0100 0001 0000 0000 0000 0100

Tag = 0x410

Set index = 0

Cache miss, set 0 is full, so tag 0x10 is replaced with 0x410 since it is last in the LRU-order list

Cache after:

0x0410004

0

miss

0x410,0x110,0x300,0x0

{ 0: {0x410, 0x110, 0x300, 0x0} }

Exam 2022-2023 question 1.1.3

Cache before: { 0: {0x410, 0x110, 0x300, 0x0} }

Reference: 0x0A00008 = 0000 1010 0000 0000 0000 0000 1000

Tag = 0xA00 Set index = 0

Cache miss, set 0 is full, so tag 0x0 is replaced with 0xA00 since it is last in the LRU-order list

Cache after:

{ 0: {0xA00, 0x410, 0x110, 0x300} }

0x0A00008	0	miss	0xA00,0x410,0x110,0x300
-----------	---	------	-------------------------

Exam 2022-2023 question 1.1.3: own example

Cache before: { 0: {0xA00, 0x410, 0x110, 0x300} }

Reference: 0xFFFFFFFF = 1111 1111 1111 1111 1111 1111 1111 1111

Tag = 0xFFFFF Set index = 127

Cache miss, but now in a new set: 127. We add the block in the new set:

Cache after:

{0: {0xA00, 0x410, 0x110, 0x300},

127: {0xFFFFF} }

Remember: only write accessed set in LRU-order!

0xFFFFFFFF	127	Miss	0xFFFFF
------------	-----	------	---------

Exam 2022-2023 question 1.1.3: own example

Cache before:

{ 0: {0xA00, 0x410, 0x110, 0x300}, 127: {0xFFFFFFFF} }

Reference: 0x4206924B = 0100 0010 0000 0110 1001 0010 0100 1011

Tag = 0x42069

Set index = 18

Cache miss, new set 18, we add the block in the new set:

Cache after:

Remember: only write accessed set in LRU-order!

{0: {0xA00, 0x410, 0x110, 0x300},

18: {0x42069}

127: {0xFFFFFFFF} }

0x42069	18	Miss	0x42069
---------	----	------	---------

Product placement!

Some friends of mine, Phillip and Mahmood, are doing a Project in Practice, and they are creating learning tools for Compsys-related topics.

So if you have run out of Compsys exam sets, then you can try out the tools:

<https://abdsecondhand.site/CACHE/dist/index.html> for some cache exercises

<https://abdsecondhand.site/VMAT/dist/index.html> for virtual table exercises

Good luck with the exam!